

1 Red-Black-Tree

1.1 Introduction and general properties

A red-black-tree is a self-balancing binary search tree. In addition to the invariants of the latter, the former guarantees that the tree remains roughly balanced even when it receives input sequences that lead to pathological cases for normal binary search trees (sorted input). Specifically, the red-black-tree guarantees that the lengths of any two paths from the root to a leaf (i.e. the depths of any two leaf nodes) differ by at most a factor of 2.

$$\forall \text{ leaf nodes } l_i, l_j \in T : \text{depth}(l_i) \leq 2 \cdot \text{depth}(l_j)$$

It does so by assigning one of two colors (traditionally red and black) to every node and maintaining the following invariants:

- (1) Every node is either red or black.
- (2) If a node is red, it doesn't have a red child.
- (3) All paths from the root to a leaf go through the same number of black nodes.

If b is the number of black nodes on every path from the root to a leaf, then the shortest such path (which contains no red nodes at all) has a length of b , and the longest such path (where black and red nodes alternate) has a length of $2b$.)

In the following discussion, let a *group* of nodes, g , in a red-black-tree denote either *any black node together with all its red child nodes*, or the *null-group* which contains no nodes at all¹.

The *weight* of a group is the number of *red* nodes it contains, obviously $0 \leq \text{weight}(g) \leq 2$. If $\text{weight}(g) = 0$, g is said to be *empty*. If $\text{weight}(g) = 2$, g is said to be *full*.

If $\text{weight}(g) = 1$, it is possible to rotate g in such a way that the former red child becomes the black root of the group, and the former black root becomes a red child on the opposite site; this rotation does not affect any descendant groups.

By invariant (2) we know that the nearest descendant nodes below a group must be black nodes, and as such they form the roots of groups themselves. Thus we can rephrase invariant (3) as follows:

- (3) All paths from the root to a leaf go through the same number of groups.

In order to preserve this invariant, inserting a node into a red-black-tree must not add a new group to the tree (except at the very root where the addition affects all paths equally). Likewise, deleting a node from the tree must not remove a group entirely (again, except at the root). The fact that groups can

¹Calling nothing a group simplifies a few arguments.

have variable numbers of red nodes provides us with enough flexibility though: We can shift nodes between groups, rotate sub-trees, and recolor nodes to make sure we have some leeway at the position where we need to insert or delete a node.

1.2 Insertion

Insertion into a normal binary search tree is straightforward. You start at the root, and then you always go left if the new value is less than the current value, or you go right if it is greater. When you've nowhere left to go (i.e. you'd need to go left but there's no left child, or the other way around) you simply add the value at that position.

Insertion into a red-black-tree works similarly but it must not add a new group somewhere down the tree because that would violate invariant (3). Therefore insertion is not possible if the direct parent of the to-be-added node is part of a full group. The solution is to transform the group the parent belongs to in such a way that it contains at most one red node. Then we can add the new node as a red child to the group, after rotating it first if necessary. The following algorithm performs this operation on the group.

```

procedure decrease-weight(n):
  p := parent of n
  pp := grandparent of n
  l, r := children of n

  // assume that n is black and l, r are red

  case 1: n is the root of the tree
    recolor:
      l, r: black

  case 2: p is black
    recolor:
      n: red
      l, r: black

  case 3: p is red
    case 3.1: pp has 1 red child (p)
      case 3.1a: pp->p->n is right-right or left-left
        right-right: rotate-left(pp)
        left-left: rotate-right(pp)
        recolor:
          pp: red
          p: black
          n: red
          l, r: black
      case 3.1b: pp->p->n is right-left or left-right
        right-left:

```

```

        rotate-right(p)
        rotate-left(pp)
    left-right:
        rotate-left(p)
        rotate-right(pp)
    recolor:
        pp: red
        l, r: black

    case 3.2: pp has 2 red children
        decrease-weight(pp)
        decrease-weight(n) // try again

```

The actual insertion algorithm makes use of this.

```

procedure insert(n, v):
    p := parent of n

    if v = n.value:
        return
    else if v < n.value and n has a left child:
        insert(n.left, v)
        return
    else if v > n.value and n has a right child:
        insert(n.right, v)
        return

    case 1: n is black
        if v < n.value:
            n.left = new node(v, red)
            n.left.parent = n
        else:
            n.right = new node(v, red)
            n.right.parent = n

    case 2: n is red
        case 2.1: weight(p) = 2
            decrease-weight(p)
            // now n can be anywhere, go again:
            insert(n, v)

        case 2.2: weight(p) = 1
            case 2.2.1: v < n.value and n = p.left
                rotate-right(p)
                n.left = new node(v, red)
            case 2.2.2: v > n.value and n = p.right
                rotate-left(p)
                n.right = new node(v, red)
            case 2.2.3: v < n.value and n = p.right
                p.left = new node(v, red)

```

```
        swap(p, p.left)
    case 2.2.4: v > n.value and n = p.left
        p.right = new node(v, red)
        swap(p, p.right)
```

1.3 Deletion

Deletion is only possible if the group we delete from is not empty.